

# A Study of Embeddings, LLMs, and RAG Methods

## Information Retrieval – Project Milestone 1

Patrascu Adrian    Modiga Miriam    Toderian Vitalii

Faculty of Automatic Control and Computer Science  
Politehnica University of Bucharest

November 2025

# Project Overview

**Goal:** Benchmark RAG retrieval methods for Information Retrieval

## Stage 1 – Implemented:

- Production RAG application
- pgvector & Qdrant stores
- Dense embeddings baseline
- NeMo Guardrails integration

## Stage 2/3 – Planned:

- ColBERT vs FAISS comparison
- Sparse embeddings (BM25, SPLADE)
- Benchmark evaluation
- Performance analysis

**Approach:** Build working baseline, then add comparison methods

# Current Implementation – Stage 1

## System Architecture:

- **Frontend:** React 19 + Vite + TailwindCSS
- **Backend:** FastAPI + LangChain
- **LLM Serving:** vLLM (Llama-3.2-3B)
- **Vector Stores:** pgvector / Qdrant
- **Embeddings:** all-MiniLM-L6-v2 (384d)

## Key Features:

- RAG toggle (enable/disable)
- PDF document upload
- Streaming chat interface
- Multiple chunking strategies
- NeMo Guardrails agents
- Dataset registry

**Deployment:** Docker Compose with GPU/CPU profiles

# ColBERT – Late Interaction Retrieval [Stage 2]

**Key Concept:** Token-level matching

**How it works:**

- ① Encode query & document with BERT
- ② Store per-token embeddings
- ③ Compute MaxSim for relevance

**Advantages:**

- High accuracy
- Pre-computed doc embeddings
- Fine-grained semantic matching

MaxSim Score

$$S_{q,d} = \sum_i \max_j E_{q_i} \cdot E_{d_j}^T$$

**Paper:** Khattab & Zaharia, SIGIR 2020

**Current:** Using standard dense embeddings with cosine similarity

# FAISS – Similarity Search at Scale [Stage 2]

**Key Concept:** Approximate nearest neighbor

**Index Types:**

- **Flat:** Exact search
- **IVF:** Inverted file
- **HNSW:** Graph-based
- **PQ:** Compressed

**Advantages:**

- Very fast queries
- GPU acceleration
- Billion-scale support
- Memory efficient

**Paper:** Johnson et al., 2017

**Current:** Using pgvector (PostgreSQL) & Qdrant for vector storage

# Embedders & LLMs

## Current Implementation:

- **Embedder:** all-MiniLM-L6-v2
  - 384 dimensions
  - ~50ms CPU, ~10ms GPU
- **LLM:** Llama-3.2-3B-Instruct
  - Served via vLLM
  - Streaming responses
- Multi-embedder architecture ready

## Planned Additions [Stage 2/3]:

- **Sparse:** BM25, SPLADE
- **Dense:** bge-base, e5-base
- **LLMs:** Mistral, Phi-3
- **Benchmarks:**
  - MS MARCO
  - Natural Questions

# Technology Stack – Current Implementation

Component	Tech
Backend	FastAPI + unicorn
Frontend	React 19 + Vite
Styling	TailwindCSS
LLM Serving	vLLM
RAG	LangChain
Vectors	pgvector, Qdrant
Guardrails	NeMo 0.13.0
Database	PostgreSQL 16

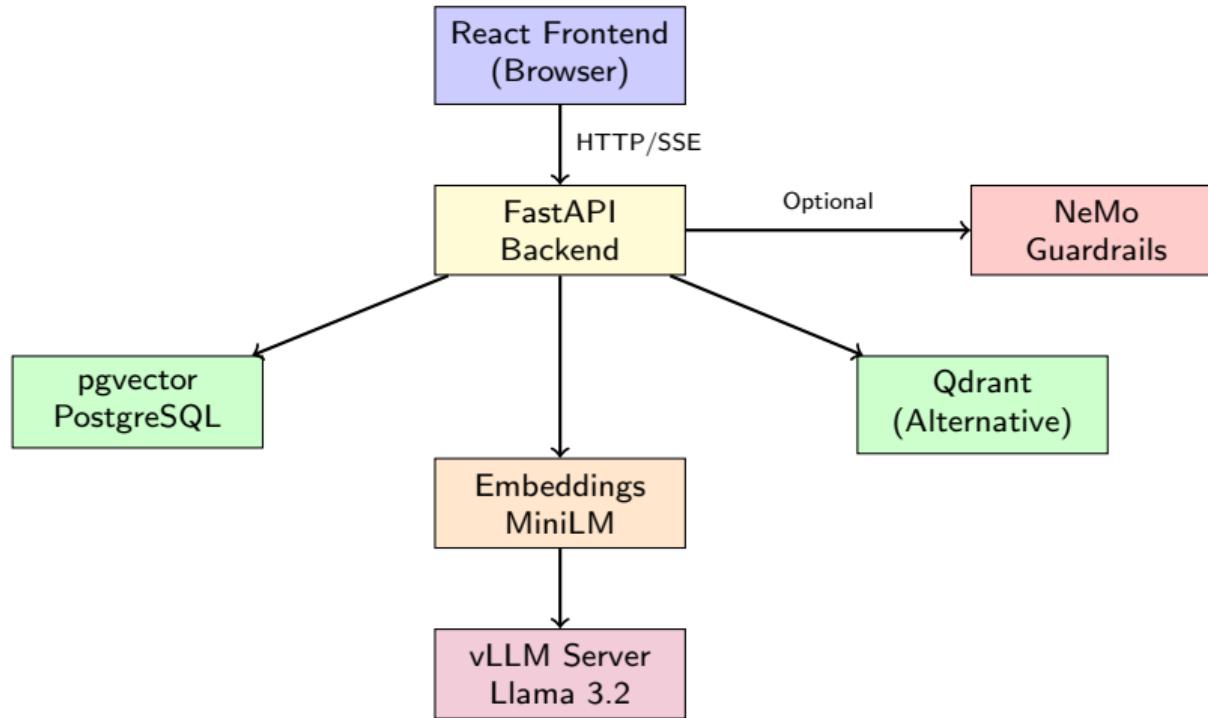
## Key Libraries:

- sentence-transformers
- PyPDF (document processing)
- Server-Sent Events (SSE)

## Deployment:

- Docker Compose
- GPU & CPU profiles
- Package manager: uv

# System Architecture – Current Implementation



**Flow:** User query → Embed → Vector search → RAG context → LLM → Stream response

# Key Features & RAG Implementation

## User-Facing Features:

- Chat interface with streaming
- RAG toggle (enable/disable)
- PDF document upload
- Multiple chat sessions
- Dark/light theme support
- NeMo Guardrails agents

## RAG Components:

- **Chunking:**
  - Recursive (default)
  - Fixed-size
  - Semantic
- **Chunk size:** 1000 chars
- **Overlap:** 200 chars (20%)
- **Retrieval:** Top-k=3, cosine similarity
- Dataset registry
- Vector store migration tools

# Evaluation Metrics [Stage 3]

## Retrieval Metrics (Planned):

- MRR@k (Mean Reciprocal Rank)
- Recall@k
- NDCG@k

## Generation Metrics (Planned):

- BLEU, ROUGE-L
- BERTScore

## Efficiency Metrics (Planned):

- Query latency (ms)
- Index size (GB)
- Memory usage
- Throughput (QPS)

## Current Testing:

- Qualitative evaluation
- Manual testing

# Challenges – Encountered & Expected

## Stage 1 Challenges (Solved):

- Docker orchestration complexity
- vLLM GPU memory optimization
- PostgreSQL pgvector setup
- Frontend-backend SSE streaming
- NeMo Guardrails integration

## Future Challenges (Stage 2/3):

- ColBERT index size management
- Fair comparison methodology
- Benchmark dataset preparation
- Hyperparameter tuning
- Evaluation automation

Aspect	Current (pgvector)	Planned (ColBERT)
Accuracy	Baseline	Higher (expected)
Speed	Fast	Slower
Index Size	Moderate	Larger

# Project Timeline & Questions

## Project Stages:

### ① Stage 1 – COMPLETED (Nov 2025):

- Production RAG baseline application
- pgvector & Qdrant integration
- NeMo Guardrails implementation
- Full-stack deployment (React + FastAPI + vLLM)

② Stage 2 – Planned: Add ColBERT, FAISS, sparse embeddings

③ Stage 3 – Planned: Benchmark evaluation (MS MARCO, NQ)

④ Stage 4 – Final: Comparative analysis & report

## Questions?

Thank you for your attention!