

A Study of Embeddings, LLMs, and RAG Methods

Information Retrieval – Project Milestone 1

Patrascu Adrian Modiga Miriam Toderian Vitalii

Faculty of Automatic Control and Computer Science
Politehnica University of Bucharest

November 2025

Project Overview

Goal: Benchmark RAG retrieval methods for Information Retrieval

Stage 1 – Implemented:

- Production RAG application
- pgvector & Qdrant stores
- Dense embeddings baseline
- NeMo Guardrails integration

Stage 2/3 – Planned:

- ColBERT vs FAISS comparison
- Sparse embeddings (BM25, SPLADE)
- Benchmark evaluation
- Performance analysis

Approach: Build working baseline, then add comparison methods

Current Implementation – Stage 1

System Architecture:

- **Frontend:** React 19 + Vite + TailwindCSS
- **Backend:** FastAPI + LangChain
- **LLM Serving:** vLLM (Llama-3.2-3B)
- **Vector Stores:** pgvector / Qdrant
- **Embeddings:** all-MiniLM-L6-v2 (384d)

Key Features:

- RAG toggle (enable/disable)
- PDF document upload
- Streaming chat interface
- Multiple chunking strategies
- NeMo Guardrails agents
- Dataset registry

Deployment: Docker Compose with GPU/CPU profiles

ColBERT – Late Interaction Retrieval [Stage 2]

Key Concept: Token-level matching

How it works:

- ① Encode query & document with BERT
- ② Store per-token embeddings
- ③ Compute MaxSim for relevance

Advantages:

- High accuracy
- Pre-computed doc embeddings
- Fine-grained semantic matching

MaxSim Score

$$S_{q,d} = \sum_i \max_j E_{q_i} \cdot E_{d_j}^T$$

Paper: Khattab & Zaharia, SIGIR 2020

Current: Using standard dense embeddings with cosine similarity

FAISS – Similarity Search at Scale [Stage 2]

Key Concept: Approximate nearest neighbor

Index Types:

- **Flat:** Exact search
- **IVF:** Inverted file
- **HNSW:** Graph-based
- **PQ:** Compressed

Advantages:

- Very fast queries
- GPU acceleration
- Billion-scale support
- Memory efficient

Paper: Johnson et al., 2017

Current: Using pgvector (PostgreSQL) & Qdrant for vector storage

Embedders & LLMs

Current Implementation:

- **Embedder:** all-MiniLM-L6-v2
 - 384 dimensions
 - ~50ms CPU, ~10ms GPU
- **LLM:** Llama-3.2-3B-Instruct
 - Served via vLLM
 - Streaming responses
- Multi-embedder architecture ready

Planned Additions [Stage 2/3]:

- **Sparse:** BM25, SPLADE
- **Dense:** bge-base, e5-base
- **LLMs:** Mistral, Phi-3
- **Evaluation Datasets:**
 - Natural Questions (NQ)
 - MS MARCO

Technology Stack – Current Implementation

Component	Tech
Backend	FastAPI + unicorn
Frontend	React 19 + Vite
Styling	TailwindCSS
LLM Serving	vLLM
RAG	LangChain
Vectors	pgvector, Qdrant
Guardrails	NeMo 0.13.0
Database	PostgreSQL 16

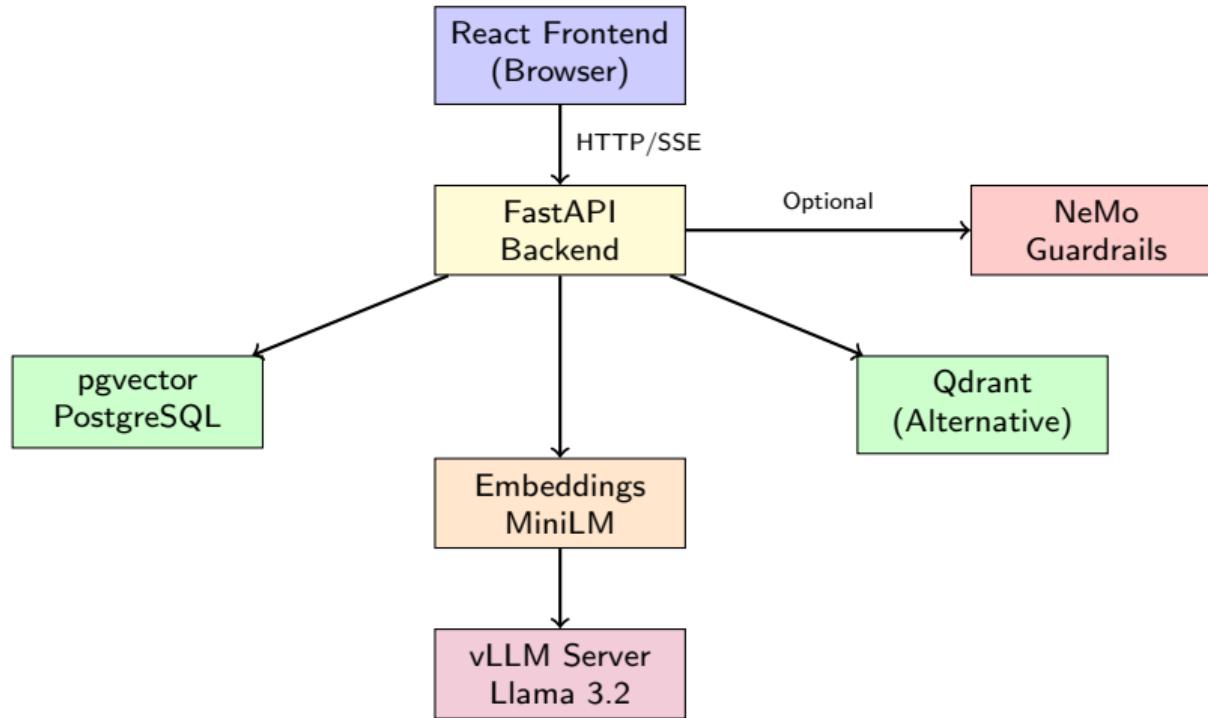
Key Libraries:

- sentence-transformers
- PyPDF (document processing)
- Server-Sent Events (SSE)

Deployment:

- Docker Compose
- GPU & CPU profiles
- Package manager: uv

System Architecture – Current Implementation



Flow: User query → Embed → Vector search → RAG context → LLM → Stream response

Key Features & RAG Implementation

User-Facing Features:

- Chat interface with streaming
- RAG toggle (enable/disable)
- PDF document upload
- Multiple chat sessions
- Dark/light theme support
- NeMo Guardrails agents

RAG Components:

- **Chunking:**
 - Recursive (default)
 - Fixed-size
 - Semantic
- **Chunk size:** 1000 chars
- **Overlap:** 200 chars (20%)
- **Retrieval:** Top-k=3, cosine similarity
- Dataset registry
- Vector store migration tools

Evaluation Datasets & Metrics [Stage 3]

Evaluation Datasets:

- **Natural Questions (NQ)**
 - Google queries + Wikipedia
 - Factoid QA (PDF use case)
- **MS MARCO**
 - Bing queries + 8.8M passages
 - Large-scale ranking

Current Testing:

- Qualitative evaluation
- Manual testing

Retrieval Metrics:

- **MRR@10**: Position of first relevant result
- **Recall@20**: % of relevant docs found
- **nDCG@10**: Ranking quality (best on top)

Generation Metrics:

- **BERTScore**: Semantic similarity with gold answer
- **Answer Relevance**: Does answer address query?

Efficiency:

- Query latency, memory, throughput

Challenges – Encountered & Expected

Stage 1 Challenges (Solved):

- Docker orchestration complexity
- vLLM GPU memory optimization
- PostgreSQL pgvector setup
- Frontend-backend SSE streaming
- NeMo Guardrails integration

Future Challenges (Stage 2/3):

- ColBERT index size management
- Fair comparison methodology
- Benchmark dataset preparation
- Hyperparameter tuning
- Evaluation automation

Aspect	Current (pgvector)	Planned (ColBERT)
Accuracy	Baseline	Higher (expected)
Speed	Fast	Slower
Index Size	Moderate	Larger

Project Timeline & Questions

Project Stages:

① Stage 1 – COMPLETED (Nov 2025):

- Production RAG baseline application
- pgvector & Qdrant integration
- NeMo Guardrails implementation
- Full-stack deployment (React + FastAPI + vLLM)

② Stage 2 – Planned: Fully Integrate ColBERT, HippoRAG and Benchmark evaluation (MS MARCO, NQ) on everything so far

③ Stage 3 – Final: Comparative analysis

Questions?

Thank you for your attention!