

# A Study of Embeddings, LLMs, and RAG Methods

Information Retrieval – Project Stage 1

Patrascu Adrian<sup>1</sup>      Modiga Miriam<sup>1</sup>  
Toderian Vitalii<sup>1</sup>

<sup>1</sup>Faculty of Automatic Control and Computer Science  
Politehnica University of Bucharest

November 2024

## 1 Project Description

This project aims to benchmark and compare two prominent Retrieval-Augmented Generation (RAG) methods: **ColBERT** and **FAISS**. RAG systems enhance Large Language Model responses by retrieving relevant documents from a knowledge base before generating answers.

### 1.1 Objectives

- Run 2-3 Information Retrieval benchmarks
- Compare 3-5 sparse vs dense embedding methods
- Evaluate 3-5 open-source LLMs for generation
- Implement conversational tests using LangGraph
- Analyze trade-offs between accuracy, speed, and memory usage

### 1.2 Scope

We will evaluate:

- **Retrieval Methods:** ColBERT (late interaction) vs FAISS (approximate nearest neighbor)
- **Embedders:** BM25, SPLADE (sparse); MiniLM, BGE, E5 (dense)
- **LLMs:** Llama 2, Mistral 7B, Phi-2, Zephyr 7B
- **Benchmarks:** MS MARCO, Natural Questions

## 2 State-of-the-Art

### 2.1 ColBERT – Late Interaction Retrieval

ColBERT [1] introduces a "late interaction" architecture that independently encodes queries and documents using BERT, then computes relevance via MaxSim operation:

$$S_{q,d} = \sum_i \max_j E_{q_i} \cdot E_{d_j}^T \quad (1)$$

**Key advantages:** Pre-computed document embeddings, token-level matching, scalable to large collections.

**Implementations:**

- Official: <https://github.com/stanford-futuredata/ColBERT>
- RAGatouille: <https://github.com/AnswerDotAI/RAGatouille> (used in this project)

## 2.2 FAISS – Similarity Search at Scale

FAISS [2] is Facebook’s library for efficient similarity search in high-dimensional spaces. It supports multiple index types:

- **Flat**: Exact search (baseline)
- **IVF**: Inverted file index for faster search
- **HNSW**: Graph-based approximate search
- **PQ**: Product quantization for compression

**Implementation:** <https://github.com/facebookresearch/faiss>

## 2.3 Related Work

- **DPR** [3]: Dense Passage Retrieval for open-domain QA
- **SPLADE** [4]: Sparse lexical and expansion model
- **BEIR** [5]: Benchmark for zero-shot IR evaluation
- **LangChain/LangGraph**: Frameworks for building LLM applications

## 3 Technologies

### 3.1 Core Stack

Component	Technology
Language	Python 3.10+
Deep Learning	PyTorch 2.0+
LLM Framework	Transformers, vLLM
RAG Pipeline	LangChain, LangGraph
ColBERT	RAGatouille (colbert-ir/colbertv2.0)
Vector Search	FAISS (CPU/GPU), PGVector
Embeddings	Sentence-Transformers
Evaluation	datasets, ranx

### 3.2 Models

**Embedders:**

- Sparse: BM25, SPLADE
- Dense: all-MiniLM-L6-v2, bge-base-en, e5-base-v2

**LLMs:** Llama 2 7B, Mistral 7B, Phi-2, Zephyr 7B

## 4 System Architecture

### 4.1 Pipeline Components

1. **Ingestion**: Load documents, chunk text, generate embeddings
2. **Indexing**: Build ColBERT index or FAISS index
3. **Retrieval**: Search for relevant passages given query
4. **Generation**: Use LLM to generate answer from context
5. **Evaluation**: Compute metrics (MRR, Recall, NDCG)




Figure 1: RAG System Architecture

## 4.2 LangGraph Integration

LangGraph orchestrates the conversational RAG pipeline:

- State management across conversation turns
- Conditional routing between retrievers
- History-aware query reformulation

## 4.3 Implementation Status

### 4.3.1 Current Progress

As of Milestone 1, the following components have been implemented on the `miriam/colbert-integration` branch:

- **ColBERT Retriever Module:** Complete implementation using RAGatouille library with the official ColBERTv2.0 model (`colbert-ir/colbertv2.0`)
- **Core Functionality:** Document indexing with late interaction search, returning ranked results with relevance scores
- **Test Suite:** Validation scripts demonstrating indexing and retrieval on sample datasets
- **Vector Database:** PGVector backend for FAISS integration
- **RAG Pipeline:** LangChain-based retrieval-augmented generation with vLLM

### 4.3.2 Technical Implementation Details

#### ColBERT Configuration:

- Model: `colbert-ir/colbertv2.0`
- Library: RAGatouille (Python wrapper)

- Max document length: 512 tokens (configurable)
- Index storage: Local filesystem (`.ragatouille` directory)
- LangChain integration: Full compatibility with Document objects

#### **Code Organization:**

- `backend/colbert_retriever.py`: Core retriever implementation (101 lines)
- `backend/test_colbert.py`: Integration testing
- `backend/rag_components.py`: RAG pipeline orchestration
- `backend/multi_embedder_manager.py`: Multi-model embedding support

#### **4.3.3 Ongoing Development**

The following features are currently under development:

- Integration of ColBERT retriever with main FastAPI application
- API endpoints for ColBERT document indexing and search
- Reranking functionality for hybrid retrieval approaches
- Comparative benchmarking framework between ColBERT and FAISS
- LangGraph integration for conversational RAG with multiple retrieval strategies

## **5 Potential Challenges**

### **5.1 Technical Challenges**

1. **Memory constraints:** ColBERT indexes can be large (storing per-token embeddings). Solution: Use compression, quantization.
2. **GPU requirements:** Running multiple LLMs requires significant VRAM. Solution: Use quantized models (4-bit), CPU offloading.
3. **Indexing time:** Building ColBERT indexes is slower than FAISS. Solution: Pre-build indexes, use batch processing.
4. **Benchmark consistency:** Ensuring fair comparison across different methods. Solution: Standardized evaluation scripts, same hardware.

### **5.2 Methodological Challenges**

1. **Hyperparameter tuning:** Each method has different optimal settings (chunk size, top-k, temperature). Solution: Grid search on validation set.
2. **Metric selection:** Different metrics favor different systems. Solution: Report multiple metrics (MRR, Recall, latency).
3. **LLM variability:** Generation quality varies with prompts. Solution: Use consistent prompt templates.

### **5.3 Expected Trade-offs**

## **6 Timeline and Deliverables**

### **6.1 Project Stages**

1. **Stage 1** (Current): Project description, architecture, technology selection
2. **Stage 2:** Implementation of retrieval pipelines

Aspect	ColBERT	FAISS
Retrieval accuracy	Higher	Lower
Query latency	Higher	Lower
Index size	Larger	Smaller
Setup complexity	More complex	Simpler

Table 1: Expected Trade-offs

3. **Stage 3:** Benchmark evaluation and results analysis
4. **Final:** Complete report with findings and recommendations

## 6.2 Milestone 1 Achievements

- **Architecture Design:** Complete system architecture for dual-retriever RAG system
- **Technology Stack:** Selection and validation of core technologies (RAGatouille, FAISS, vLLM, LangChain)
- **ColBERT Implementation:** Working ColBERT retriever module with indexing and search capabilities
- **Development Environment:** Dockerized setup with PostgreSQL, vLLM, and frontend
- **Test Framework:** Unit tests and integration tests for retrieval components

## 6.3 Next Steps for Stage 2

- Merge ColBERT integration branch into main codebase
- Implement unified API endpoints for both retrieval methods
- Develop comparative benchmarking scripts using MS MARCO dataset
- Create evaluation metrics dashboard
- Integrate multiple embedding models (BM25, SPLADE, MiniLM, BGE, E5)
- Conduct initial performance testing on target hardware

## References

- [1] O. Khattab and M. Zaharia, “Colbert: Efficient and effective passage search via contextualized late interaction over bert,” in *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 39–48, ACM, 2020.
- [2] J. Johnson, M. Douze, and H. Jégou, “Billion-scale similarity search with gpus,” *IEEE Transactions on Big Data*, vol. 7, no. 3, pp. 535–547, 2019.
- [3] V. Karpukhin, B. Oğuz, S. Min, P. Lewis, L. Wu, S. Edunov, D. Chen, and W.-t. Yih, “Dense passage retrieval for open-domain question answering,” in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 6769–6781, 2020.
- [4] T. Formal, B. Piwowarski, and S. Clinchant, “Splade: Sparse lexical and expansion model for first stage ranking,” in *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 2288–2292, ACM, 2021.
- [5] N. Thakur, N. Reimers, A. Rücklé, A. Srivastava, and I. Gurevych, “Beir: A heterogeneous benchmark for zero-shot evaluation of information retrieval models,” in *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2021.