

A Study of Embeddings, LLMs, and RAG Methods

Information Retrieval – Project Stage 1

Patrascu Adrian¹ Modiga Miriam¹
Toderian Vitalii¹

¹Faculty of Automatic Control and Computer Science
Politehnica University of Bucharest

November 2024

1 Project Description

This project aims to benchmark and compare two prominent Retrieval-Augmented Generation (RAG) methods: **ColBERT** and **FAISS**. RAG systems enhance Large Language Model responses by retrieving relevant documents from a knowledge base before generating answers.

1.1 Objectives

- Run 2-3 Information Retrieval benchmarks
- Compare 3-5 sparse vs dense embedding methods
- Evaluate 3-5 open-source LLMs for generation
- Implement conversational tests using LangGraph
- Analyze trade-offs between accuracy, speed, and memory usage

1.2 Scope

We will evaluate:

- **Retrieval Methods:** ColBERT (late interaction) vs FAISS (approximate nearest neighbor)
- **Embedders:** BM25, SPLADE (sparse); MiniLM, BGE, E5 (dense)
- **LLMs:** Llama 2, Mistral 7B, Phi-2, Zephyr 7B
- **Benchmarks:** MS MARCO, Natural Questions

2 State-of-the-Art

2.1 ColBERT – Late Interaction Retrieval

ColBERT [1] introduces a "late interaction" architecture that independently encodes queries and documents using BERT, then computes relevance via MaxSim operation:

$$S_{q,d} = \sum_i \max_j E_{q_i} \cdot E_{d_j}^T \quad (1)$$

Key advantages: Pre-computed document embeddings, token-level matching, scalable to large collections.

Implementations:

- Official: <https://github.com/stanford-futuredata/ColBERT>
- RAGatouille: <https://github.com/AnswerDotAI/RAGatouille>

2.2 FAISS – Similarity Search at Scale

FAISS [2] is Facebook’s library for efficient similarity search in high-dimensional spaces. It supports multiple index types:

- **Flat**: Exact search (baseline)
- **IVF**: Inverted file index for faster search
- **HNSW**: Graph-based approximate search
- **PQ**: Product quantization for compression

Implementation: <https://github.com/facebookresearch/faiss>

2.3 Related Work

- **DPR** [3]: Dense Passage Retrieval for open-domain QA
- **SPLADE** [4]: Sparse lexical and expansion model
- **BEIR** [5]: Benchmark for zero-shot IR evaluation
- **LangChain/LangGraph**: Frameworks for building LLM applications

3 Technologies

3.1 Core Stack

Component	Technology
Language	Python 3.10+
Deep Learning	PyTorch 2.0+
LLM Framework	Transformers, vLLM
RAG Pipeline	LangChain, LangGraph
ColBERT	RAGatouille
Vector Search	FAISS (CPU/GPU)
Embeddings	Sentence-Transformers
Evaluation	datasets, ranx

3.2 Models

Embedders:

- Sparse: BM25, SPLADE
- Dense: all-MiniLM-L6-v2, bge-base-en, e5-base-v2

LLMs: Llama 2 7B, Mistral 7B, Phi-2, Zephyr 7B

4 System Architecture

4.1 Pipeline Components

1. **Ingestion**: Load documents, chunk text, generate embeddings
2. **Indexing**: Build ColBERT index or FAISS index
3. **Retrieval**: Search for relevant passages given query
4. **Generation**: Use LLM to generate answer from context
5. **Evaluation**: Compute metrics (MRR, Recall, NDCG)




Figure 1: RAG System Architecture

4.2 LangGraph Integration

LangGraph orchestrates the conversational RAG pipeline:

- State management across conversation turns
- Conditional routing between retrievers
- History-aware query reformulation

5 Potential Challenges

5.1 Technical Challenges

1. **Memory constraints:** ColBERT indexes can be large (storing per-token embeddings). Solution: Use compression, quantization.
2. **GPU requirements:** Running multiple LLMs requires significant VRAM. Solution: Use quantized models (4-bit), CPU offloading.
3. **Indexing time:** Building ColBERT indexes is slower than FAISS. Solution: Pre-build indexes, use batch processing.
4. **Benchmark consistency:** Ensuring fair comparison across different methods. Solution: Standardized evaluation scripts, same hardware.

5.2 Methodological Challenges

1. **Hyperparameter tuning:** Each method has different optimal settings (chunk size, top-k, temperature). Solution: Grid search on validation set.

- Metric selection:** Different metrics favor different systems. Solution: Report multiple metrics (MRR, Recall, latency).
- LLM variability:** Generation quality varies with prompts. Solution: Use consistent prompt templates.

5.3 Expected Trade-offs

Aspect	ColBERT	FAISS
Retrieval accuracy	Higher	Lower
Query latency	Higher	Lower
Index size	Larger	Smaller
Setup complexity	More complex	Simpler

Table 1: Expected Trade-offs

6 Timeline and Deliverables

- Stage 1** (Current): Project description, architecture, technology selection
- Stage 2:** Implementation of retrieval pipelines
- Stage 3:** Benchmark evaluation and results analysis
- Final:** Complete report with findings and recommendations

References

- [1] O. Khattab and M. Zaharia, “Colbert: Efficient and effective passage search via contextualized late interaction over bert,” in *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 39–48, ACM, 2020.
- [2] J. Johnson, M. Douze, and H. Jégou, “Billion-scale similarity search with gpus,” *IEEE Transactions on Big Data*, vol. 7, no. 3, pp. 535–547, 2019.
- [3] V. Karpukhin, B. Oğuz, S. Min, P. Lewis, L. Wu, S. Edunov, D. Chen, and W.-t. Yih, “Dense passage retrieval for open-domain question answering,” in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 6769–6781, 2020.
- [4] T. Formal, B. Piwowarski, and S. Clinchant, “Splade: Sparse lexical and expansion model for first stage ranking,” in *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pp. 2288–2292, ACM, 2021.
- [5] N. Thakur, N. Reimers, A. Rücklé, A. Srivastava, and I. Gurevych, “Beir: A heterogeneous benchmark for zero-shot evaluation of information retrieval models,” in *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2021.